

Enriching Historic Photography with Structured Data using Image Region Segmentation

Taylor Arnold[†], Lauren Tilton^{*}

[†]Department of Mathematics and Computer Science, University of Richmond

^{*}Department of Rhetoric and Communication Studies, University of Richmond

410 Westhampton Way, Richmond, VA, USA 23173

{tarnold2, ltilton}@richmond.edu

Abstract

Modern computer vision algorithms provide methods for automatically splitting the content of an image into labelled regions. While early segmentation methods focused on detecting regions corresponding to concrete objects, recent work has considered the detection of unenumerable areas of “stuff” such as the sky, water, and roads. In this paper we describe a method for using detected stuff regions, augmented with the location of people, to produce rich and accurate structured data triples for describing the content of historic photography. We apply our method to a collection of 1600 documentary photographs produced in the 1930s and 1940 by the FSA-OWI division of the U.S. federal government. Manual verification of the extracted annotations yields an accuracy rate of 97.5%, compared to 70.7% for relations extracted from object detection and 31.5% for automatically generated captions. Our method also produces a rich set of features, providing more unique labels (1170) than either the captions (1040) or thing category (178) methods. We conclude by describing directions for more linguistically-focused ontology of region categories that can better enrich historical image data. Open source code and the extracted metadata from our corpus are made available as external resources.

Keywords: computer vision, image segmentation, cultural heritage, photography, Linked Data, ontology, digital humanities

1. Introduction

Museums, archives, libraries, and other cultural heritage organizations have increasingly looked to provide structured metadata about historic collections in an effort to increase access and discovery. Where records have been digitized and rights restrictions allow for it, many of these organizations have also been able to make the digital records directly accessible through openly available APIs and URIs. Prominent examples of these efforts include the Rijksmuseum’s *RijksData* (Dijkshoorn et al., 2018), Europeana’s Search API, Record API, and SPARQL endpoint (Concordia et al., 2009), and the *Linked Data Service* provided by the United States Library of Congress (Zimmer, 2015). The effort to make resources available within a cohesive semantic web offers exciting possibilities for research and public access to cultural collections. Yet, challenges remain in fully realizing the potential of structured data for opening up digital archives.

Many digital collections held by cultural heritage organizations consist of still and moving image data. These include scans of textual documents, photographs of material culture, and digital scans of newspapers, artwork, photographs and other visual objects. Unlike machine-readable textual archives, visual collections do not immediately offer a simple method for automated search or data extraction. While records often include extensive metadata about the provenance of a digital image, there is often no structured data pertaining to the content of the image itself. Even when descriptive captions exist, these are often short and intended to be read alongside the object itself. Captions are often written assuming that the reader will be able to view the object itself. The lack of structured linguistic descriptions serves as a roadblock to providing rich links between and

across collections, as well as limiting the possibilities for large-scale analysis of collections. While expert and crowd-sourced annotations can fill in some gaps, manual data construction requires extensive resources and is not capable of keeping up with the increasing size of digitized datasets (Seitsonen, 2017).

Computer vision techniques provide a direction for the automated production of structured data to enrich collections of historic digital images. Machine learning techniques can detect features present in images and store these alongside human-generated metadata pertaining to the digital records. However, the use of automated techniques have their own unique set of challenges. Most computer vision algorithms are built using modern datasets, and may produce annotations that are inaccurate or inappropriate for historic data. Incorrectly extracted data records are particularly concerning when making data available to the public. Even when including confidence scores for extracted features, studies have shown that people have trouble accurately interpreting probabilistic data and are overall confidence in predictions (Khaw et al., 2019). The challenges of mis-classified data are particularly acute when they risk reinforcing racial, gender, and socioeconomic biases inherent in the training of machine learning techniques. For example, a recent study showed that face detection algorithms have difficulty identifying darker skinned individuals (Buolamwini and Gebru, 2018). Applying state-of-the-art face detection algorithms to a collection of photographs, therefore, risks further hiding marginalized communities.

In this article we present a method for the automated extraction of highly-accurate structured data describing the content of historic photography using computer vision algorithms. Specifically, our approach is based on the detection of regions of the image containing elements described as




Figure 1: Eight selected object types from the MS COCO dataset, each shown with the five images from the FSA-OWI color photographs that are most predicted to contain the given category (Wu et al., 2019). All categories were estimated to exist a probability greater than 80%. The plane and horse categories seem to have correctly identified the objects in their five respective images, and two of the cow images are in fact cows (the others are horses). The other categories seem to be all false detections. Some mistakes can be easily explained, such as the color of the plane in the first row resembling a banana or the large grey statues looking like elephants. However, many mistakes are hard to explain, such as the row of skateboard objects.



Figure 2: Three detected captions for three FSA-OWI photographs using the ‘Show, attend and tell’ model (Xu et al., 2015). The first provides a caption that matches the image and the third produces a caption that is very similar to the image. The second correctly identifies the subject of a woman in the frame but mistakenly believes she is holding a microphone. The final image produces an annotation that incorrectly labels the people as giraffes.

“stuff”, which includes elements such as sky, water, trees, grass, and roads (Caesar et al., 2018). We use an algorithm to detect which regions of an image contain these features and produced structured data. While temporal, cultural, and regional differences can exist in some of these categories, the stuff-based regions of images are significantly more robust than many other features that can currently be extracted from image data.

The remainder of this article is structured as follows. Section 2. gives a brief survey of several projects currently using computer vision and structured data to augment historic image collections. In Section 3. we provide an overview of image segmentation and the current approaches for the classification of stuff categories. Section 4. presents our specific approach and schema for producing structured data from images and Section 5. presents an evaluation of our approach applied to a collection of 1600 photographs from the 1930s and 1940s. We conclude in Section 6. with a discussion of future possibilities and challenges of applying image segmentation to historic datasets.

2. Background

2.1. Object Detection

The identification of objects within an image constitutes one of the most popular and important tasks in computer vision. Early tasks focused on relatively simple objectives, such as the classification of hand-written digits in the MNIST dataset, which used small 28-by-28 black and white pixels (Platt, 1999). Modern training datasets feature thousands of categories, ranging from very specific concrete items such as specific species of dogs to relatively abstract concepts such as ‘grocery stores’ and ‘parties’. Using transfer learning, in which a model trained on one dataset is modified to function on a new task, it is possible to produce algorithms trained to detect virtually any object category by manually tagging only a small set of training examples. The training of models for specific features has been employed in the annotation of several historical image datasets, such as the location of Dadaism art work (Thompson and Mimno, 2017) and detecting figures in digitized newspapers (Wevers and Smits, 2019).

Object detection is a useful tool for annotating specific features of interest within a collection. However, each feature requires a manually trained model and may not generalize well to a new collection. Using existing models with pre-selected categories on historic images typically produces a mix of correct annotations and false annotations. Figure 1 shows the results of applying Detectron2’s object detection algorithms to the FSA-OWI collection (Wu et al., 2019). While some categories produced reasonably accurate annotations, such as the detection of horses and people, most categories detected more false positives than successfully generated tags. Without a good general-purpose collection of object detectors, a challenge discussed further in Section 6., makes object detection difficult to use as a means for producing structured data for linking historic image collections.

2.2. Automated Captions

The automated generation of descriptive image captions is a more ambitious task that has been a popular line of research at the intersection of computational linguistics and computer vision. Captions generated with the help of linked textual data have shown to be fairly accurate, offering a useful tool for automated description of images in news articles and other media powerful (Hessel et al., 2019) (Batra et al., 2018) (Hollink et al., 2016). As with object detection, automatically generated captions within well-defined domains, such as profile photos, has also been fairly successful at generating accurate descriptions (Gatt et al., 2018). On the more general task of generating free-form image captions, current state-of-the-art methods also produce impressive results when applied to modern datasets (Nikolaus et al., 2019) (Jiang et al., 2019a) (Wang et al., 2018). On datasets that differ from the specific training data, however, modern methods too-often produce nonsensical results that make it difficult to use their results directly in an archive. Figure 2 shows the results of one popular caption algorithm applied to photographs from the 1940s (Xu et al., 2015). While two captions produce reasonable results, a third incorrectly identifies the object held by the main subject and the fourth mistakenly believes the two men in the frame are giraffes.

2.3. Image Embedding

Given the difficulty of automatically producing accurate structure data from image collections, the use of image embedding has become a popular approach for finding links between and across collections of visual data. Similar to the process of using word embeddings, image embeddings most frequently project a collection of images into the penultimate layer of a supervised neural network. Once represented as a sequence of numbers in a high-dimensional space, images within an across collections can be associated with their closest neighbors (McAuley et al., 2015). Flattening image embeddings into two or three dimensions produces useful visualizations of large image collections. Tools in the digital humanities, such as Yale DH Lab's *Pix-Plot*, make this approach accessible to a large community of users and illustrates the appeal of its method (Duhaime, 2019).

For finding similar images or detecting patterns and trends within a collection, image embeddings are a useful tool and generalize well to new and historic datasets. By forgoing the explicit creation of structured data, they avoid many of the pitfalls of the automated information extraction. However, the constructed data does not produce meaningful relationships that can be easily distributed as structured data. This makes it difficult to extend the recommendation system to new collections and to find links across a web of archives.

2.4. Ontologies for Visual Data

The utility of structured data rests on describing data using standard ontologies. It is important, when extracting data for linkage and discovery, to carefully consider the schema(s) to use in describing relationships. There currently exist several ontologies for describing image data. Schema.org supplies generic schemas for photographs, images, paintings, and creative works (Guha et al., 2016). Dublin core offers a well-established ontology for describing digital records specifically designed for libraries and digital archives (Weibel, 1997). Both of these are useful for describing the provenance of digital objects. Several schema also exist for describing the content of image data, often with a specific focus on describing time-coded moving images such as film and television. The *Advene* project provides an ontology designed to integrate with their manually annotation tool (Aubert and Prié, 2005). The *Audio-Visual Rhetorics of Affect* group extended this vocabulary to include more granular terms that capture formal elements of affect and film studies (Agt-Rickauer et al., 2018).

The field-specific ontologies provided for digital images provide useful methods for linking collections. Our digital project based on the FSA-OWI collection uses the Dublin Core Metadata Element Set to describe each record. In our work here, however, we aim for simplicity by describing our annotations using a class extension of the the Web Annotation Data Model (Sanderson et al., 2017).

3. Image segmentation

Current state-of-the-art models for detecting objects within images are difficult to use as a general-purpose code system for the analysis of visual culture. Currently available models feature categories are too specific and only cover a very small number of the object types that could be seen within the frame of modern, western-centric film and photography. When considering historical or more diverse datasets, the coverage is even worse. For example, the COCO dataset contains only three types of fruits (apple, orange, banana), two vegetables (broccoli, carrot), and five other food items (cake, donut, pizza, hotdog, sandwich). There are no generic catch-all food categories for other items falling outside of these lists. Applying these object detection models indiscriminately to a large corpus without understanding its limitations will result in biased results. They will find certain kinds of food items, animals, and clothing, but will completely ignore examples outside of a narrowly curated list of categories.

In 2018, a research team from University of Edinburgh and Google AI released a newly hand-tagged version of the COCO dataset that contained 91 new categories. The dominant focus of computer vision research at that time had been the identification and localization of objects. This, they argued, had overshadowed the automatic identification other categories that constitute the "amorphous background regions" within an image (Caesar et al., 2018). They described the existing categories within the COCO dataset as consisting only of "thing" categories. That is, these categories consist of objects that have a specific size, well-defined shape, concrete parts, and can be enumerated. Regions that do not correspond to things, such as the sky, water, and the ceiling, often compose the majority of an image but at the time they lacked the data and models needed for automatically identifying them. The team described these regions, in contrast to things, as "stuff". Their work resulted in a comprehensive ontology for "stuff" regions. Their approach proposes to split all regions under two super-categories: "indoor stuff" and "outdoor stuff". These groups are further divided into mid-level categories, which include "water", "ground", "sky", "furniture", and "floor". Finally, these are split into 91 fine-grained categories such as "sea", "mud", "clouds", and "carpet". The team labelled the entirety of the COCO dataset with these classifications. The joint task of identifying these labels alongside the "thing" labels, collectively known as the panoptic task, has been one of the primary competitions sponsored by the COCO dataset in 2018 and 2019 (Kirillov et al., 2019). As a result, there are now many accurate models for automatically labelling these regions within an image. Table 1 provides the hierarchy of categories used in the 2019 challenge and Figure 3 shows an example of the segmentation algorithm applied to an image from the FSA-OWI archive.

While no classification scheme can be free of cultural assumptions nor account for all possible scenarios, the stuff categories are significantly more generic than the object categories. This is particularly true of the high- and mid-level categories. The higher-level categories avoid some of the material specific designations from the lowest-level

Group	Metacategories	Categories
indoor	building	bridge; building; house; roof; tent
indoor	ceiling	ceiling
indoor	floor	floor; floor-wood
indoor	food	food; fruit
indoor	furniture	cabinet; counter; door-stuff; light; mirror-stuff; shelf; stairs; table
indoor	rawmaterial	cardboard; paper
indoor	textile	banner; blanket; curtain; pillow; rug; towel
indoor	wall	wall; wall-brick; wall-stone; wall-tile; wall-wood
indoor	window	window; window-blind
outdoor	ground	dirt; gravel; pavement; platform; playingfield; railroad; road; sand; snow
outdoor	plant	flower; grass; tree
outdoor	sky	sky
outdoor	solid	mountain; rock
outdoor	structural	fence; net
outdoor	water	river; sea; water

Table 1: Hierarchical description of stuff categories, modified from the dataset produced for the MS COCO Panopticon task.

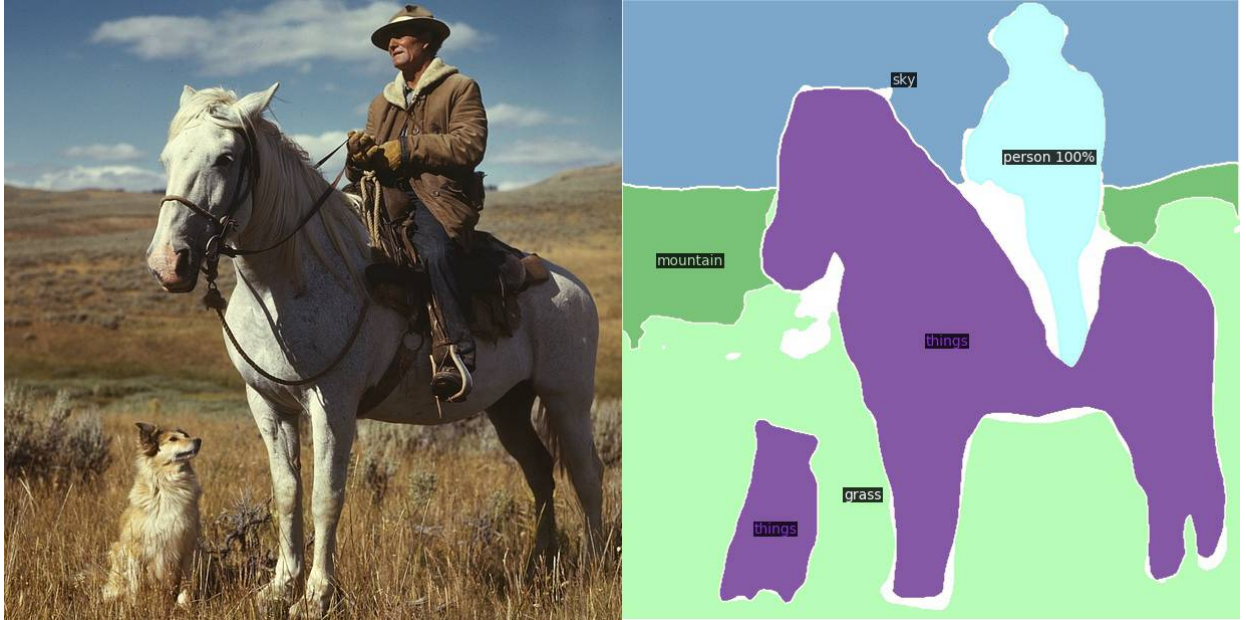


Figure 3: Example of the Panoptic FPN (R101-FPN) stuff-segmentation algorithm applied to a region of one FSA-OWI photograph (Wu et al., 2019). The algorithm detected five types of regions: sky, mountain, grass, things, and person.

categories, such as wood-flooring, that may not be applicable with images that significantly depart from the available training data. By aggregating information about detected stuff categories, we can make intelligent guesses about whether an image was taken inside or outside, how the people in the image are placed relative to the background, and the location and role of the horizon in framing the image.

As always when working with automatically generated annotations, care should be taken to avoid misinterpreting the results of stuff-segmentation algorithms. There are some categories that have some ambiguity between them, such as “dirt” and “sand” or “mat” and “rug”. Also, the stuff categories were designed pragmatically for the task of assigning all of the pixels in an image to a fixed set of classi-

fications. The distinction between stuff and things is not a sharp epistemological distinction. Several categories overlap between the two, such as “furniture” and “door”; the difference in labels is a result of the size of the images and their resolution rather than a fundamental property of the objects themselves. These ambiguities are essentially unavoidable and should not deter the usage of the stuff categories. The only caution is to avoid making claims that may come down to relatively arbitrary distinctions between categories—for example, claiming that Photographer A took more photos with dirt backgrounds whereas Photographer B preferred sand backgrounds—without carefully evaluating the appropriateness of the distinction and the accuracy of the automatic identification in a particular application.

```

@prefix rdf:      <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix pgram:    <http://photogrammar.org#> .
@prefix dcterms:  <http://purl.org/dc/terms/> .
@prefix oa:       <http://www.w3.org/ns/oa#> .

<http://photogrammar.org/anno1> a oa:Annotation ;
  dcterms:creator <http://photogrammar.org/tarnold2> ;
  dcterms:created "2020-02-19T12:00:00Z" ;
  oa:hasBody [
    a                pgram:ImageSegmentationRegion ;
    pgram:regionName <http://example.org/stuff/things> ;
    pgram:regionPercent 32 ;
  ] ;
  oa:hasTarget <http://photogrammar.org/resource/1a35022v> ;
  oa:motivatedBy oa:tagging .

<http://photogrammar.org/anno2> a oa:Annotation ;
  dcterms:creator <http://photogrammar.org/tarnold2> ;
  dcterms:created "2020-02-19T12:01:00Z" ;
  oa:hasBody [
    a                pgram:ImageSegmentationRegion ;
    pgram:regionName <http://example.org/stuff/people> ;
    pgram:regionPercent 6 ;
    pgram:regionCount 1 ;
  ] ;
  oa:hasTarget <http://photogrammar.org/resource/1a35022v> ;
  oa:motivatedBy oa:tagging .

```

Schema 1: Example of extracted structured data from the image in Figure 3 using the stuff-based image segmentation technique.

4. Method

Our proposed method for the automatic extraction of structured data from image data begins by applying the Detectron2 implementation of the Panoptic FPN (R101-FPN) model to an image for stuff and thing segmentation (Wu et al., 2019). The total proportion of the image allocated to each stuff category is computed from the annotated image. We then store the percentage of each image that falls under a specific category, for any category that constitutes more than 5% of the total image. Additionally, we also tabulate the number of detected people in the image. While the general COCO object detections are not reliable on historic images, the detection of the people category is reasonably accurate across different corpora and the presence (or absence) of people within an image is an important feature to distinguish different image subjects.

Schema 1 shows any example of the extracted structured data from regions detected in the image from Figure 3. Each detected region type within an image is assigned a unique identifier describing the region. This region is then associated with the original image, the type of region and the percentage of area taken up by the region. For the person annotation, the number of individual objects (1) is also recorded. Not shown in the example is a structured description of the region type codes that encode the hierarchical relationships described in Table 1. The title of the image is included to indicate where other image-level metadata would be recorded—such as the photographer, date, and rights information—in the full record.

5. Evaluation

The annotation method described in Section 4. was applied to the entire corpus of 1610 color photographs from the FSA-OWI collection (Trachtenberg, 1990). An example of these are show in Figure 4. Two additional annotations were also computed. Each photograph was tagged with detected objects from MS COCO and labelled with any object that appeared with at least a probability of 85% and for each photograph an automatically detected caption was produced (Figures 1-2 show examples of these annotations). The annotations for each photograph were coded to indicate where the annotation was accurately applied. A stuff region label was considered accurate if the region was visible within the image and a thing region label was considered accurate if the object existed somewhere in the image. A caption was considered accurate if it could be considered true in a strictly literal sense. For example, a caption saying that there are two people in an image that contains three people was considered correct for our purpose. Because not all images are guaranteed to include a region that falls above our threshold for inclusion, we also recorded the percentage of images that had at least one corresponding label (called recall in the results). The results are given in Table 2.

Both the close-analysis of the annotations in Figure 4 illustrate the efficacy of stuff region-based annotations for adding structured data to historic image data. The things annotations do offer many useful features, but have an error rate around 30%, making them difficult to deploy without




Figure 4: Seven selected stuff types and the people from the MS COCO panopticon dataset, each shown with the five images from the FSA-OWI color photographs that are most predicted to contain the given type (Wu et al., 2019). The only labels that appears the be falsely detected are in the third and fifth birdge images, where construction equipment is falsely believed to be a bridge.

	Acc.	Recall	Unique Results
Stuff & People	97.5%	98.9%	1140
Things	70.7%	37.3%	178
Captions	31.5%	100%	1040

Table 2: Results of manually validated labels produced on the FSA-OWI color photographs.

manual validation. At the moment the captions are correct less than a third of the time, and even the best captions fall far short of human-produced records. The stuff regions have an accuracy of 97.5%; while public display of estimated should contain a note about their auto-generated nature, it is possible to use these annotations without manual validation. The high accuracy of the stuff-based annotation method does not come at the cost of producing only uninteresting or unexpressive relations. In fact, the number of uniquely labelled images is slightly high than even the captions-based method, and labels were found for nearly 99% of all images. Looking manually at the results of the most representative images, we see that the stuff-categories capture key features of most of the image backgrounds and many of their foregrounds.

6. Conclusions and Future Directions

We have presented a technique for the automated production of structured data describing the content of photographic-corpora. The robustness and efficacy of our method was shown through a case-study using 1600 documentary photographs from the 1930s and 1940s. While other methods, such as object-detection and automated caption generation, have the potential to provide additional structured data, the generalizability of our approach suggests that it should be seriously considered in any attempt to algorithmically enrich large corpora of photographic materials for public dissemination.

The approach presented here offers several avenues for further extensions to supply additional structured information to historic image corpora. First, there are a number of ways that we could further encode information about the detected regions. For example, recording the dominant colors of each region type or indicating what part of an image a region is located. Secondly, it is possible to develop a structured language for creating image captions from the structured data. In connection with the first item, this would lead to captions such as a “Photograph of two people, with a green mountain and blue sky in the background”. This could produce image captions that, while more predictable than techniques allowing for free-form language, are also significantly more accurate. Finally, and most ambitiously, would be to construct a generic, hierarchical version of the MS COCO object detection algorithm that simulates the stuff-based regions. This would allow for a similar usage of object-detection algorithms for the automated extraction of objects in the foreground of an image.

7. Acknowledgements

Work supported by a National Endowment for the Humanities Digital Advancement Grant (HAA-261239-18) and Andrew Mellon Foundations *Collections as Data: Part to Whole* initiative.

8. Bibliographical References

- Agt-Rickauer, H., Hentschel, C., and Sack, H. (2018). Semantic annotation and automated extraction of audiovisual staging patterns in large-scale empirical film studies. In *SEMANTICS Posters&Demos*.
- Alexiev, V. (2018). Museum linked open data: Ontologies, datasets, projects. *Digital Presentation and Preservation of Cultural and Scientific Heritage*, (VIII):19–50.
- Alikhani, M., Nag Chowdhury, S., de Melo, G., and Stone, M. (2019). CITE: A corpus of image-text discourse relations. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 570–575, Minneapolis, Minnesota, June. Association for Computational Linguistics.
- Aubert, O. and Prié, Y. (2005). Advence: active reading through hypervideo. In *Proceedings of the sixteenth ACM conference on Hypertext and hypermedia*, pages 235–244.
- Baldwin, S. (1968). Poverty and politics; the rise and decline of the farm security administration.
- Batra, V., He, Y., and Vogiatzis, G. (2018). Neural Caption Generation for News Images. In Nicoletta Calzolari (Conference chair), et al., editors, *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 7-12, 2018. European Language Resources Association (ELRA).
- Buolamwini, J. and Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on fairness, accountability and transparency*, pages 77–91.
- Caesar, H., Uijlings, J., and Ferrari, V. (2018). Coco-stuff: Thing and stuff classes in context. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 1209–1218.
- Chen, H., Zhang, H., Chen, P.-Y., Yi, J., and Hsieh, C.-J. (2018). Attacking visual language grounding with adversarial examples: A case study on neural image captioning. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2587–2597, Melbourne, Australia, July. Association for Computational Linguistics.
- Concordia, C., Gradmann, S., and Siebinga, S. (2009). Not (just) a repository, nor (just) a digital library, nor (just) a portal: A portrait of europeana as an api. In *World Library and Information Congress: 75th IFLA General Conference and Council*.
- Derby, S., Miller, P., Murphy, B., and Devereux, B. (2018). Using sparse semantic embeddings learned from multimodal text and image data to model human conceptual

- knowledge. In *Proceedings of the 22nd Conference on Computational Natural Language Learning*, pages 260–270, Brussels, Belgium, October. Association for Computational Linguistics.
- Dijkshoorn, C., Jongma, L., Aroyo, L., Van Ossenbruggen, J., Schreiber, G., Ter Weele, W., and Wielemaker, J. (2018). The rijksmuseum collection as linked data. *Semantic Web*, 9(2):221–230.
- Duhaime, D. (2019). Pixplot: Visualize large image collections with webgl. <https://github.com/YaleDHLab/pix-plot>.
- Fan, Z., Wei, Z., Wang, S., and Huang, X. (2019). Bridging by word: Image grounded vocabulary construction for visual captioning. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 6514–6524, Florence, Italy, July. Association for Computational Linguistics.
- Gatt, A., Tanti, M., Muscat, A., Paggio, P., Farrugia, R. A., Borg, C., Camilleri, K., Rosner, M., and der Plas, L. V. (2018). Face2Text: Collecting an Annotated Image Description Corpus for the Generation of Rich Face Descriptions. In Nicoletta Calzolari (Conference chair), et al., editors, *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 7-12, 2018. European Language Resources Association (ELRA).
- Gella, S. and Keller, F. (2018). An evaluation of image-based verb prediction models against human eye-tracking data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 758–763, New Orleans, Louisiana, June. Association for Computational Linguistics.
- Guha, R. V., Brickley, D., and Macbeth, S. (2016). Schema.org: evolution of structured data on the web. *Communications of the ACM*, 59(2):44–51.
- Hartmann, M. and Søgaard, A. (2018). Limitations of cross-lingual learning from image search. In *Proceedings of The Third Workshop on Representation Learning for NLP*, pages 159–163, Melbourne, Australia, July. Association for Computational Linguistics.
- Hessel, J., Savva, N., and Wilber, M. (2015). Image representations and new domains in neural image captioning. In *Proceedings of the Fourth Workshop on Vision and Language*, pages 29–39, Lisbon, Portugal, September. Association for Computational Linguistics.
- Hessel, J., Lee, L., and Mimno, D. (2019). Unsupervised discovery of multimodal links in multi-image, multi-sentence documents. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2034–2045, Hong Kong, China, November. Association for Computational Linguistics.
- Hollink, L., Bedjeti, A., van Harmelen, M., and Elliott, D. (2016). A corpus of images and text in online news. In Nicoletta Calzolari (Conference Chair), et al., editors, *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC 2016)*, Paris, France, may. European Language Resources Association (ELRA).
- Jiang, M., Hu, J., Huang, Q., Zhang, L., Diesner, J., and Gao, J. (2019a). REO-relevance, extraneousness, omission: A fine-grained evaluation for image captioning. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 1475–1480, Hong Kong, China, November. Association for Computational Linguistics.
- Jiang, M., Huang, Q., Zhang, L., Wang, X., Zhang, P., Gan, Z., Diesner, J., and Gao, J. (2019b). TIGER: Text-to-image grounding for image caption evaluation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2141–2152, Hong Kong, China, November. Association for Computational Linguistics.
- Khaw, M. W., Stevens, L., and Woodford, M. (2019). Individual differences in the perception of probability. Available at SSRN 3446790.
- Kirillov, A., He, K., Girshick, R., Rother, C., and Dollár, P. (2019). Panoptic segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 9404–9413.
- Kiros, J., Chan, W., and Hinton, G. (2018). Illustrative language understanding: Large-scale visual grounding with image search. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 922–933, Melbourne, Australia, July. Association for Computational Linguistics.
- Lin, T.-Y., RoyChowdhury, A., and Maji, S. (2015). Bilinear CNN models for fine-grained visual recognition. In *Proceedings of the IEEE international conference on computer vision*, pages 1449–1457.
- Madhyastha, P., Wang, J., and Specia, L. (2019). VIFIDEL: Evaluating the visual fidelity of image descriptions. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 6539–6550, Florence, Italy, July. Association for Computational Linguistics.
- Mason, R. and Charniak, E. (2012). Apples to oranges: Evaluating image annotations from natural language processing systems. In *Proceedings of the 2012 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 172–181, Montréal, Canada, June. Association for Computational Linguistics.
- McAuley, J., Targett, C., Shi, Q., and Van Den Hengel, A. (2015). Image-based recommendations on styles and substitutes. In *Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 43–52.
- Mensink, T. and Van Gemert, J. (2014). The rijksmuseum challenge: Museum-centered visual recognition. In *Pro-*

- ceedings of International Conference on Multimedia Retrieval*, pages 451–454.
- Nikolaus, M., Abdou, M., Lamm, M., Aralikkat, R., and Elliott, D. (2019). Compositional generalization in image captioning. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 87–98, Hong Kong, China, November. Association for Computational Linguistics.
- Platt, J. C. (1999). Using analytic qp and sparseness to speed training of support vector machines. In *Advances in neural information processing systems*, pages 557–563.
- Sadler, P., Scheffler, T., and Schlangen, D. (2019). Can neural image captioning be controlled via forced attention? In *Proceedings of the 12th International Conference on Natural Language Generation*, pages 427–431, Tokyo, Japan, October–November. Association for Computational Linguistics.
- Sanderson, R., Ciccarese, P., and Young, B. (2017). Web annotation data model. <https://www.w3.org/TR/annotation-vocab/>. Accessed: 2020-02-19.
- Seitsonen, O. (2017). Crowdsourcing cultural heritage: public participation and conflict legacy in finland. *Journal of Community Archaeology & Heritage*, 4(2):115–130.
- Sharma, P., Ding, N., Goodman, S., and Soricut, R. (2018). Conceptual captions: A cleaned, hypernymed, image alt-text dataset for automatic image captioning. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2556–2565, Melbourne, Australia, July. Association for Computational Linguistics.
- Shimizu, N., Rong, N., and Miyazaki, T. (2018). Visual question answering dataset for bilingual image understanding: A study of cross-lingual transfer using attention maps. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 1918–1928, Santa Fe, New Mexico, USA, August. Association for Computational Linguistics.
- Singhal, K., Raman, K., and ten Cate, B. (2019). Learning multilingual word embeddings using image-text data. In *Proceedings of the Second Workshop on Shortcomings in Vision and Language*, pages 68–77, Minneapolis, Minnesota, June. Association for Computational Linguistics.
- Thompson, L. and Mimno, D. (2017). Computational cut-ups: The influence of dada. *The Journal of Modern Periodical Studies*, 8(2):179–195.
- Trachtenberg, A. (1990). *Reading American Photographs: Images as History-Matthew Brady to Walker Evans*. Macmillan, London, England.
- van Miltenburg, E., Elliott, D., and Vossen, P. (2018). Measuring the diversity of automatic image descriptions. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 1730–1741, Santa Fe, New Mexico, USA, August. Association for Computational Linguistics.
- Vinyals, O., Toshev, A., Bengio, S., and Erhan, D. (2016). Show and tell: Lessons learned from the 2015 mscoco image captioning challenge. *IEEE transactions on pattern analysis and machine intelligence*, 39(4):652–663.
- Wang, J., Madhyastha, P. S., and Specia, L. (2018). Object counts! bringing explicit detections back into image captioning. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 2180–2193, New Orleans, Louisiana, June. Association for Computational Linguistics.
- Weibel, S. (1997). The dublin core: a simple content description model for electronic resources. *Bulletin of the American Society for Information Science and Technology*, 24(1):9–11.
- Wevers, M. and Smits, T. (2019). The visual digital turn: Using neural networks to study historical images. *Digital Scholarship in the Humanities*.
- Wu, Y., Kirillov, A., Massa, F., Lo, W.-Y., and Girshick, R. (2019). Detectron2. <https://github.com/facebookresearch/detectron2>.
- Xu, K., Ba, J., Kiros, R., Cho, K., Courville, A., Salakhudinov, R., Zemel, R., and Bengio, Y. (2015). Show, attend and tell: Neural image caption generation with visual attention. In *International conference on machine learning*, pages 2048–2057.
- Yokota, M. and Nakayama, H. (2018). Augmenting Image Question Answering Dataset by Exploiting Image Captions. In Nicoletta Calzolari (Conference chair), et al., editors, *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 7-12, 2018. European Language Resources Association (ELRA).
- Zhao, S., Sharma, P., Levinboim, T., and Soricut, R. (2019). Informative image captioning with external sources of information. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 6485–6494, Florence, Italy, July. Association for Computational Linguistics.
- Zimmer, M. (2015). The twitter archive at the library of congress: Challenges for information practice and information policy. *First Monday*, 20(7).

9. Language Resource References